



# 15-721 DATABASE SYSTEMS

## Lecture #09 – Storage Models & Data Layout

---

Andy Pavlo // Carnegie Mellon University // Spring 2016

# TODAY'S AGENDA


---

In-Memory Data Layout

Storage Models

Project #2: Performance Profiling

# DATA ORGANIZATION



# DATA ORGANIZATION

---

One can think of an in-memory database as just a large array of bytes.

→ The schema tells the DBMS how to convert the bytes into the appropriate type.

Each tuple is prefixed with a header that contains its meta-data.

Storing tuples with just their fixed-length data makes it easy to compute the starting point of any tuple.

# DATA REPRESENTATION

---

## INTEGER/BIGINT/SMALLINT/TINYINT

→ C/C++ Representation

## NUMERIC

→ IEEE-754 Standard

## VARCHAR/VARBINARY/TEXT/BLOB

→ Pointer to other location if type is  $\geq 64$ -bits

→ Header with length and address to next location (if segmented), followed by data bytes.

## TIME/DATE/TIMESTAMP

→ 32/64-bit integer of (micro)seconds since Unix epoch

# DATA REPRESENTATION

---

```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  value BIGINT  
);
```

*char[]*



# DATA REPRESENTATION

---

```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  value BIGINT  
);
```



# DATA REPRESENTATION

```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  value BIGINT  
);
```





# DATA REPRESENTATION

```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  value BIGINT  
);
```



# DATA REPRESENTATION

```
CREATE TABLE JoySux (  
→ id INT PRIMARY KEY,  
  value BIGINT  
);
```



```
reinterpret_cast<int32_t*>(address)
```

# NULL DATA TYPES

---

## Choice #1: Special Values

- Designate a value to represent **NULL** for a particular data type (e.g., **INT32\_MIN**).

## Choice #2: Null Column Bitmap Header

- Store a bitmap in the tuple header that specifies what attributes are null.

## Choice #3: Per Attribute Null Flag

- Store a flag that marks that a value is null.
- Have to use more space than just a single bit because this messes up with word alignment.

# NULL DATA TYPES

## Integer Numbers

Data Type	Size	Size (Not Null)	Synonyms	Min Value	Max Value
BOOL	2 bytes	1 byte	BOOLEAN	0	1
BIT	9 bytes	8 bytes			
TINYINT	2 bytes	1 byte		-128	127
SMALLINT	4 bytes	2 bytes		-32768	32767
MEDIUMINT	4 bytes	3 bytes		-8388608	8388607
INT	8 bytes	4 bytes	INTEGER	-2147483648	2147483647
BIGINT	12 bytes	8 bytes		$-2^{63}$	$(2^{63}) - 1$

this messes up with word alignment.

# NULL DATA TYPES

## Integer Numbers

Data Type	Size	Size (Not Null)	Synonyms	Min Value	Max Value
BOOL	2 bytes	1 byte	BOOLEAN	0	1
BIT	9 bytes	8 bytes			
TINYINT	2 bytes	1 byte		-128	127
SMALLINT	4 bytes	2 bytes		-32768	32767
MEDIUMINT	4 bytes	3 bytes		-8388608	8388607
INT	8 bytes	4 bytes	INTEGER	-2147483648	2147483647
BIGINT	12 bytes	8 bytes		$-2^{63}$	$(2^{63}) - 1$

this messes up with word alignment.

# NULL DATA TYPES

---

## Choice #1: Special Values

- Designate a value to represent **NULL** for a particular data type (e.g., **INT32\_MIN**).

## Choice #2: Null Column Bitmap Header

- Store a bitmap in the tuple header that specifies what attributes are null.

## Choice #3: Per Attribute Null Flag

- Store a flag that marks that a value is null.
- Have to use more space than just a single bit because this messes up with word alignment.

# NOTICE

---

The truth is that you only need to worry about word-alignment for cache lines (e.g., 64 bytes).

I'm going to show you the basic idea using 64-bit words since it's easier to see...

# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  cdate TIMESTAMP,  
  color CHAR(2),  
  zipcode INT  
);
```

*char[]*

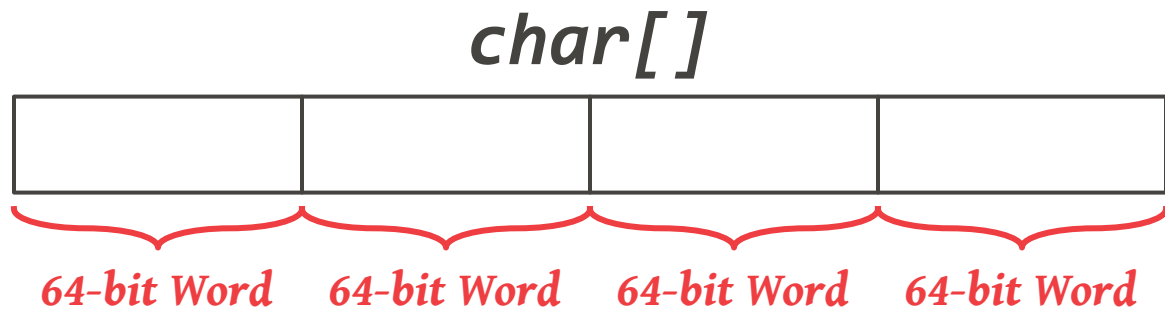




# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

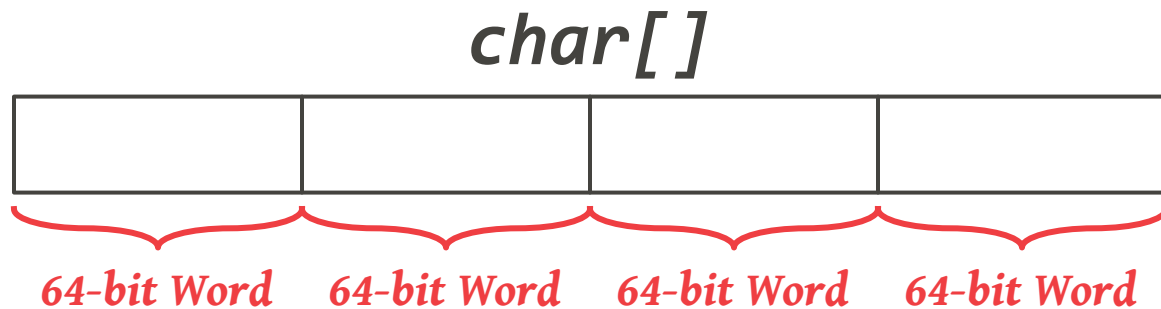
```
CREATE TABLE JoySux (  
  id INT PRIMARY KEY,  
  cdate TIMESTAMP,  
  color CHAR(2),  
  zipcode INT  
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

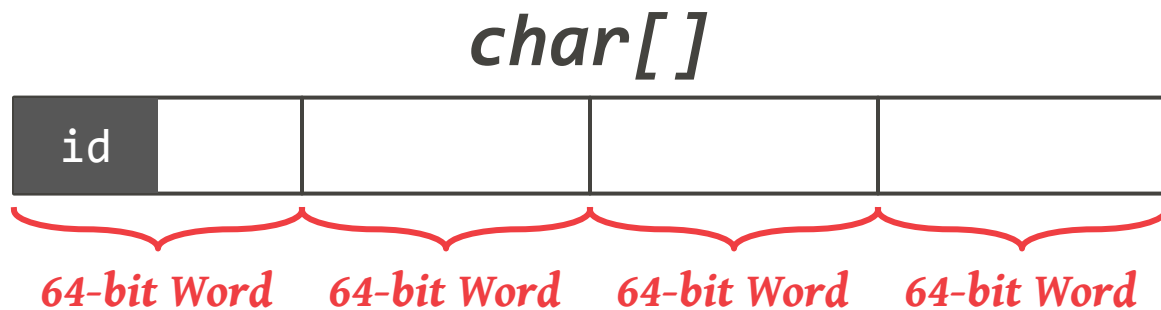
```
CREATE TABLE JoySux (  
32-bits id INT PRIMARY KEY,  
cdate TIMESTAMP,  
color CHAR(2),  
zipcode INT  
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

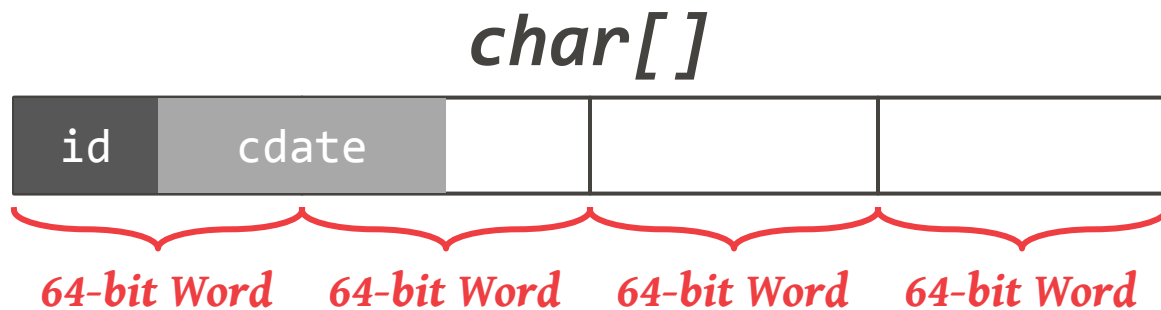
```
CREATE TABLE JoySux (  
32-bits id INT PRIMARY KEY,  
cdate TIMESTAMP,  
color CHAR(2),  
zipcode INT  
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

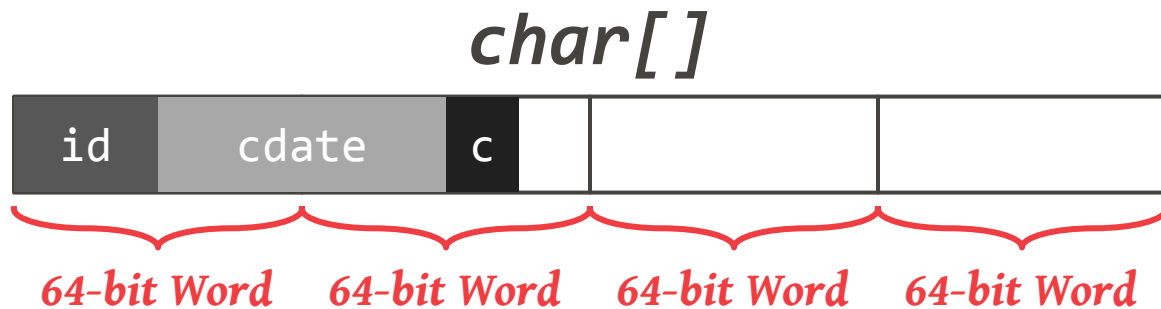
```
CREATE TABLE JoySux (  
  32-bits id INT PRIMARY KEY,  
  64-bits cdate TIMESTAMP,  
  color CHAR(2),  
  zipcode INT  
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

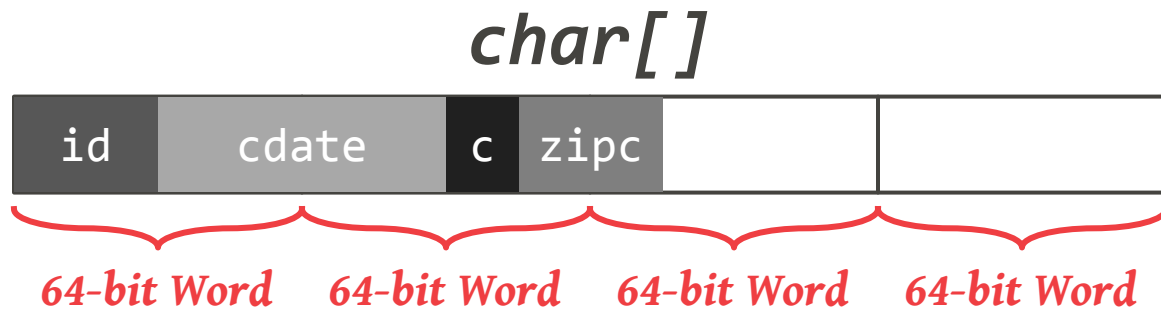
```
CREATE TABLE JoySux (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  zipcode INT
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

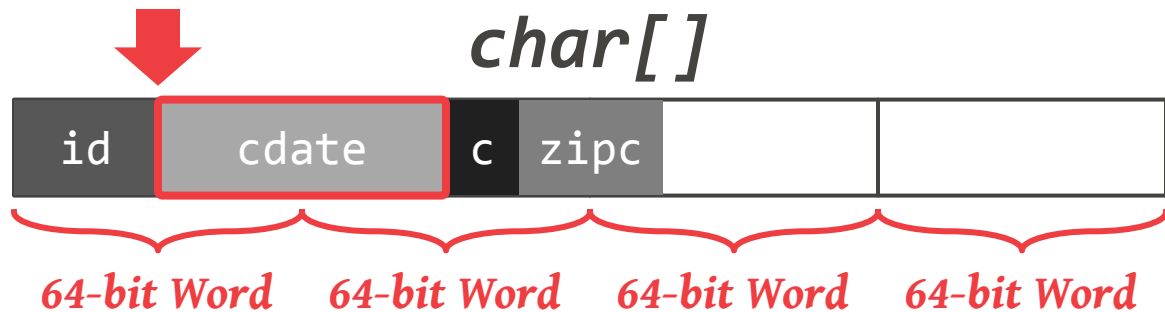
```
CREATE TABLE JoySux (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  32-bits zipcode INT
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE JoySux (
  id INT PRIMARY KEY,
  cdate TIMESTAMP,
  color CHAR(2),
  zipcode INT
);
```



# WORD-ALIGNED TUPLES

---

If the CPU fetches a 64-bit value that is not word-aligned, it has four choices:

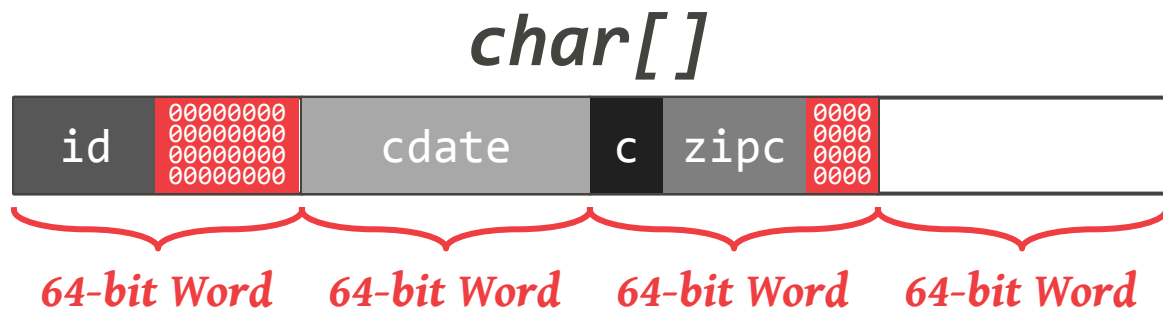
- Execute two reads to load the appropriate parts of the data word and reassemble them.
- Read some unexpected combination of bytes assembled into a 64-bit word.
- Throw an exception



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

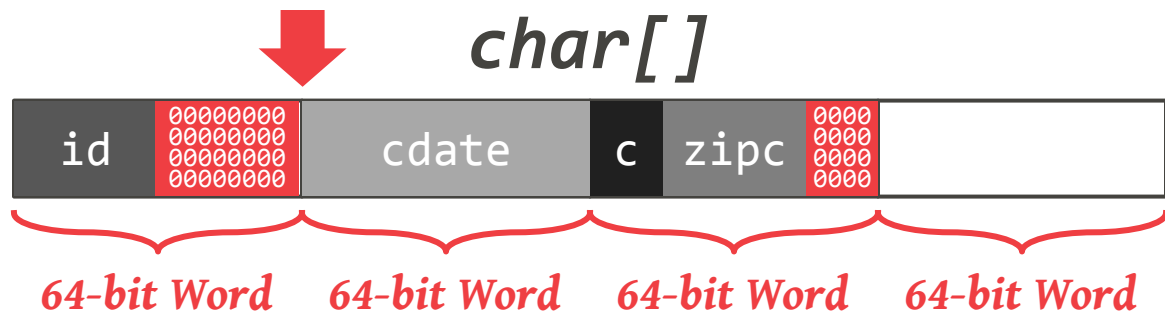
```
CREATE TABLE JoySux (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  32-bits zipcode INT
);
```



# WORD-ALIGNED TUPLES

All attributes in a tuple must be word aligned to enable the CPU to access it without any unexpected behavior or additional work.

```
CREATE TABLE JoySux (
  32-bits id INT PRIMARY KEY,
  64-bits cdate TIMESTAMP,
  16-bits color CHAR(2),
  32-bits zipcode INT
);
```



# STORAGE MODELS

---

*N*-ary Storage Model (NSM)

Decomposition Storage Model (DSM)

Hybrid Storage Model

## N-ARY STORAGE MODEL (NSM)

---

The DBMS stores all of the attributes for a single tuple contiguously.

Ideal for OLTP workloads where txns tend to operate only on an individual entity and insert-heavy workloads.

Use the tuple-at-a-time iterator model.

# NSM PHYSICAL STORAGE

---

## **Choice #1: Heap-Organized Tables**

- Tuples are stored in blocks called a heap.
- The heap does not necessarily define an order.

## **Choice #2: Index-Organized Tables**

- Tuples are stored in the index itself.
- Not quite the same as a clustered index.

# CLUSTERED INDEXES

---

The table is stored in the sort order specified by the primary key.

→ Can be either heap- or index-organized storage.

Some DBMSs always use a clustered index.

→ If a table doesn't include a pkey, the DBMS will automatically make a hidden row id pkey.

Other DBMSs cannot use them at all.

→ A clustered index is non-practical in a MVCC DBMS using the **Insert Method**.

# N-ARY STORAGE MODEL (NSM)

---

## Advantages

- Fast inserts, updates, and deletes.
- Good for queries that need the entire tuple.
- Can use index-oriented physical storage.

## Disadvantages

- Not good for scanning large portions of the table and/or a subset of the attributes.

# DECOMPOSITION STORAGE MODEL (DSM)

---

The DBMS stores a single attribute for all tuples contiguously in a block of data.

→ Sometimes also called vertical partitioning.

Ideal for OLAP workloads where read-only queries perform large scans over a subset of the table's attributes.

Use the vector-at-a-time iterator model.



# DECOMPOSITION STORAGE MODEL (DSM)

---

**1970s:** Cantor DBMS

**1980s:** DSM Proposal

**1990s:** SybaseIQ (in-memory only)

**2000s:** Vertica, Vectorwise, MonetDB

**2010s:** “The Big Three”  
Cloudera Impala, Amazon Redshift,  
SAP HANA, MemSQL

# CLUSTERED INDEXES

---

Some columnar DBMSs store data in sorted order to maximize compression.

→ Bitmap indexes with RLE from last class

Vertica does not even use indexes because all columns are sorted.

# TUPLE IDENTIFICATION

## Choice #1: Fixed-length Offsets

→ Each value is the same length for an attribute.

## Choice #2: Embedded Tuple Ids

→ Each value is stored with its tuple id in a column.

### *Offsets*

	A	B	C	D
0				
1				
2				
3				

### *Embedded Ids*

	A	B	C	D
0		0		0
1		1		1
2		2		2
3		3		3

# DECOMPOSITION STORAGE MODEL (DSM)

---

## **Advantages**

- Reduces the amount wasted work because the DBMS only reads the data that it needs.
- Better compression (last lecture).

## **Disadvantages**

- Slow for point queries, inserts, updates, and deletes because of tuple splitting/stitching.

# OBSERVATION

---

Data is “hot” when first entered into database

→ A newly inserted tuple is more likely to be updated again the near future.

As a tuple ages, it is updated less frequently.

→ At some point, a tuple is only accessed in read-only queries along with other tuples.

What if we want to use this data to make decisions that affect new txns?

# BIFURCATED ENVIRONMENT

---




*OLTP Data Silos*



*OLAP Data Warehouse*

# BIFURCATED ENVIRONMENT



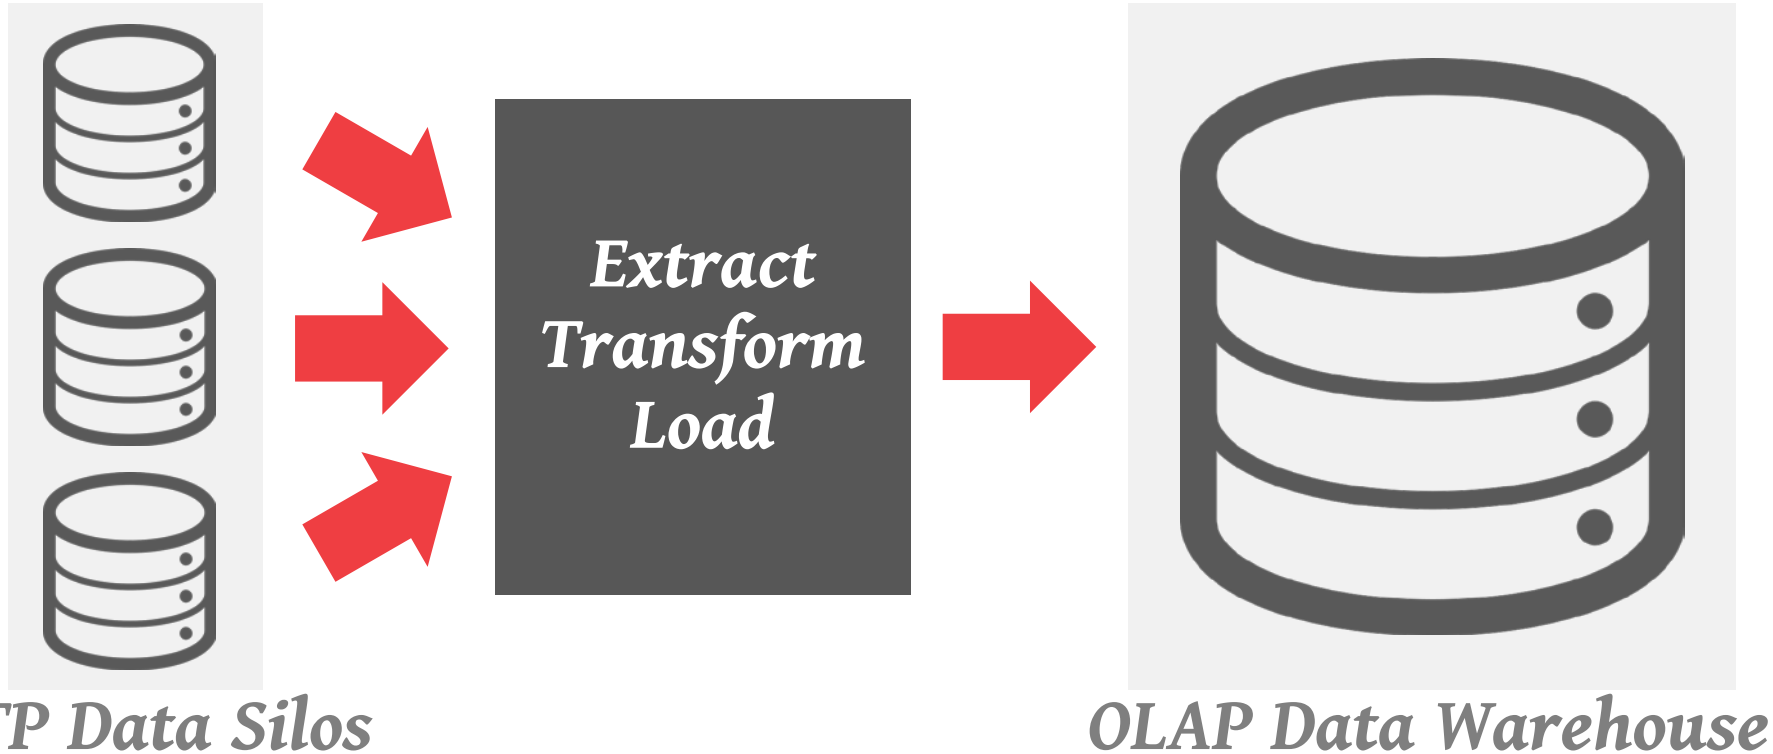
*Extract  
Transform  
Load*



*OLTP Data Silos*

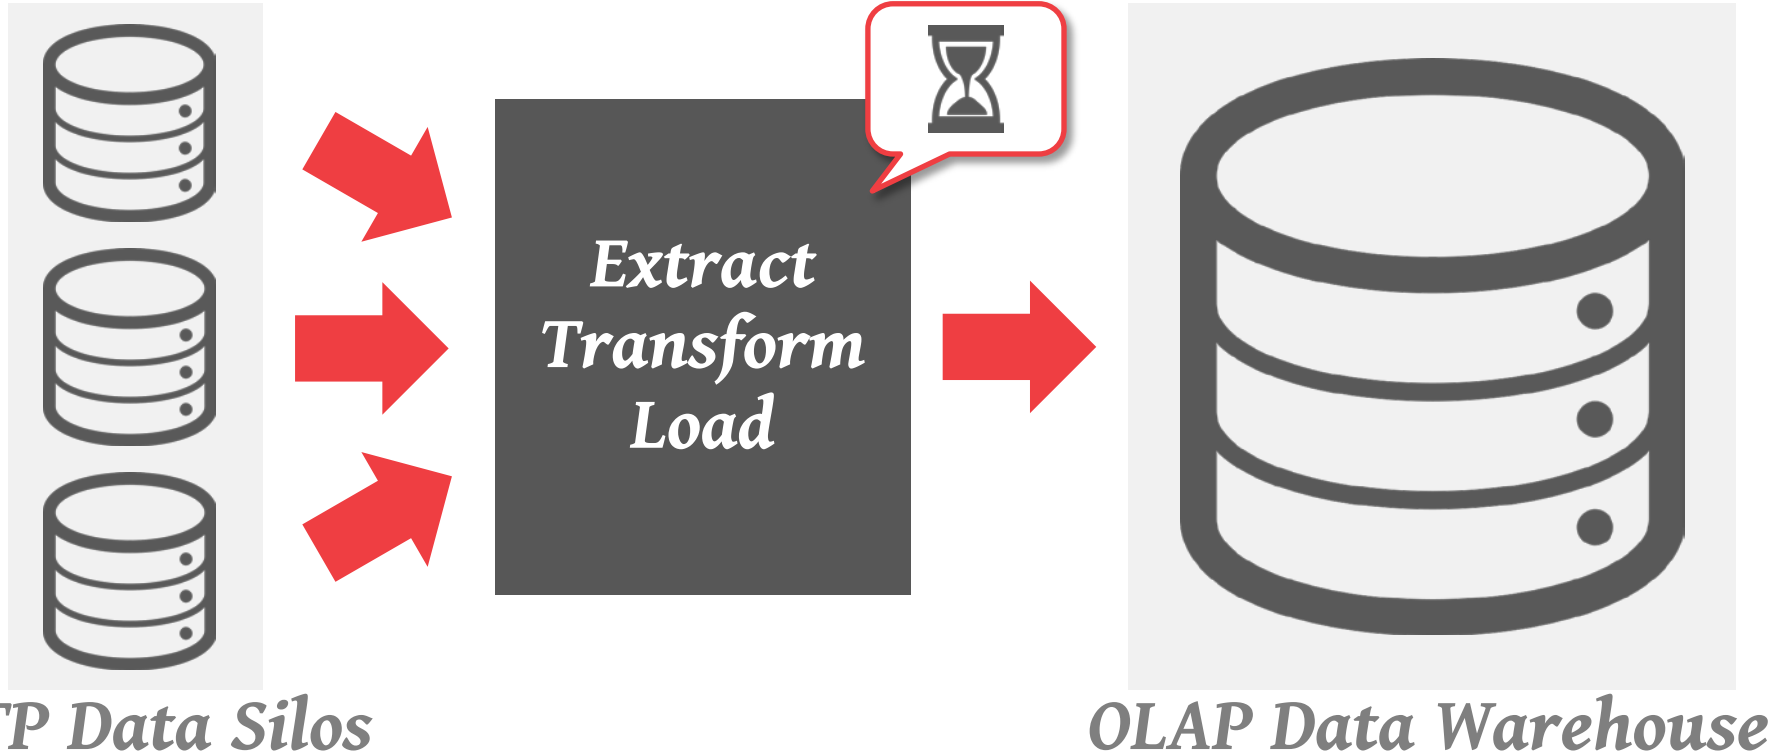
*OLAP Data Warehouse*

# BIFURCATED ENVIRONMENT

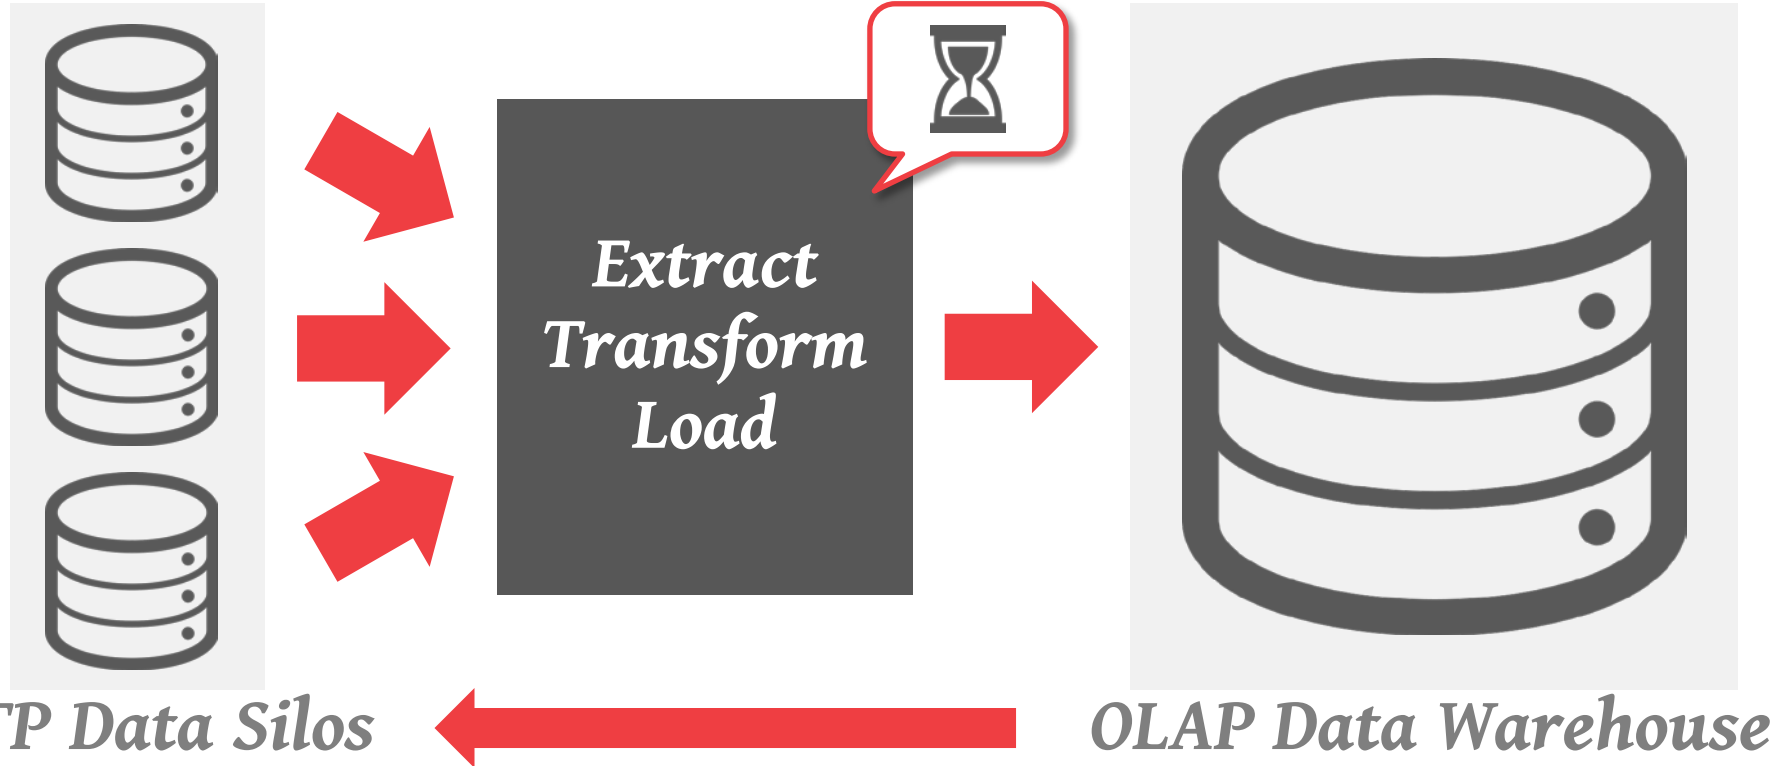




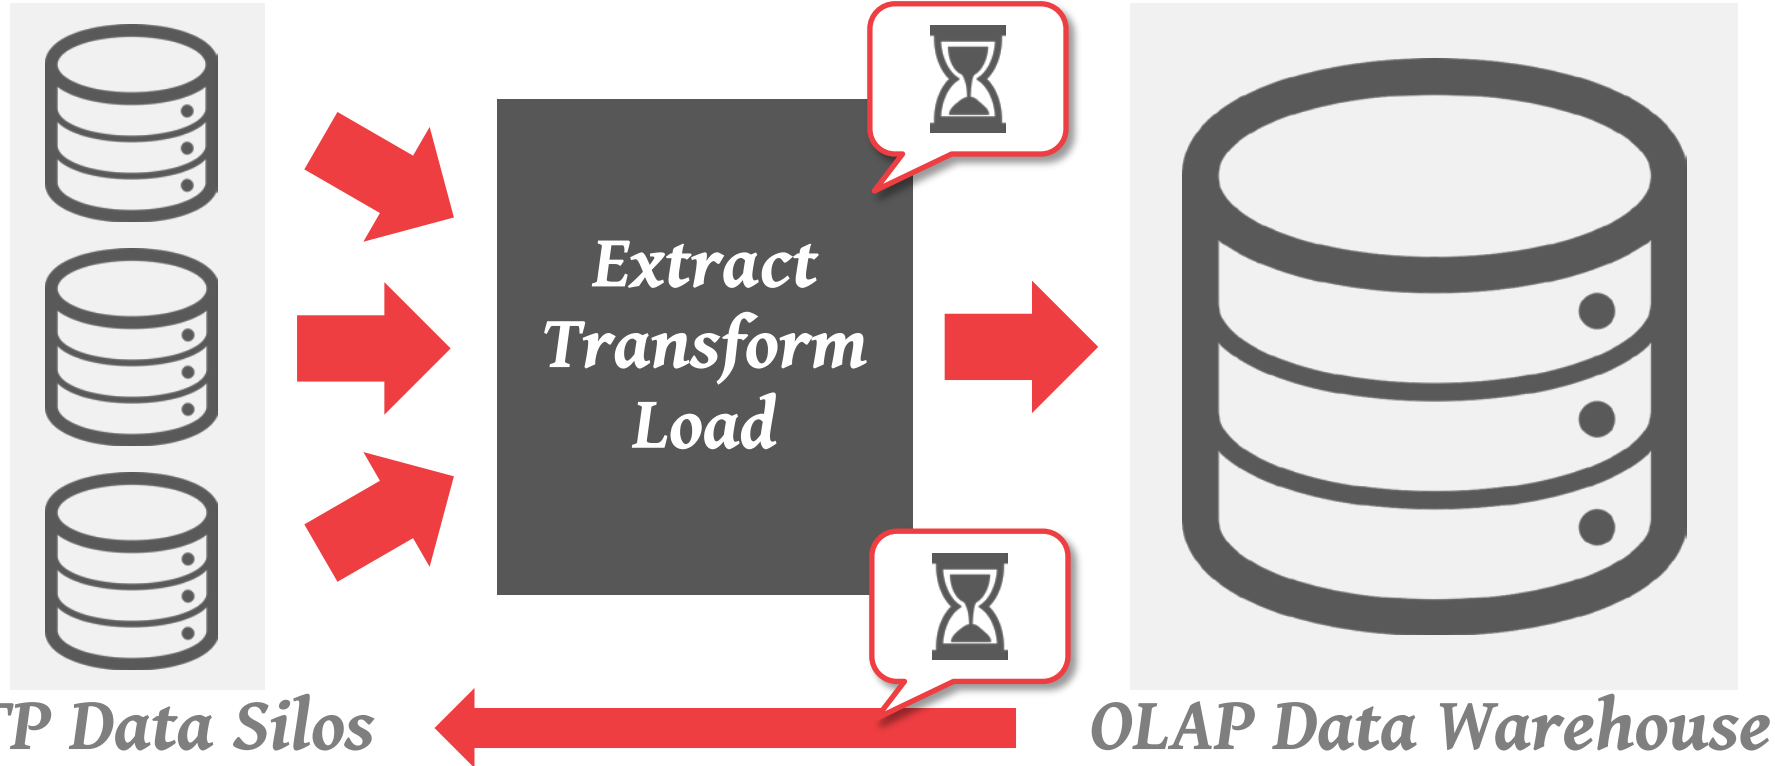
# BIFURCATED ENVIRONMENT



# BIFURCATED ENVIRONMENT



# BIFURCATED ENVIRONMENT



# HYBRID STORAGE MODEL

---

Single logical database instance that uses different storage models for hot and cold data.

Store new data in NSM for fast OLTP  
Migrate data to DSM for more efficient OLAP

# HYBRID STORAGE MODEL

---

## **Choice #1: Separate Execution Engines**

- Use separate execution engines that are optimized for either NSM or DSM databases.

## **Choice #2: Single, Flexible Architecture**

- Use single execution engine that is able to efficiently operate on both NSM and DSM databases.

# SEPARATE EXECUTION ENGINES

---

Run separate “internal” DBMSs that each only operate on DSM or NSM data.

- Need to combine query results from both engines to appear as a single logical database to the application.
- Have to use a synchronization method (e.g., 2PC) if a txn spans execution engines.

Two approaches to do this:

- Fractured Mirrors (Oracle, IBM)
- Delta Store (SAP HANA)

# FRACTURED MIRRORS

---

Store a second copy of the database in a DSM layout that is automatically updated.

→ All updates are first entered in NSM then eventually copied into DSM mirror.




A CASE FOR FRACTURED MIRRORS  
*VLDB 2002*

# FRACTURED MIRRORS

---

Store a second copy of the database in a DSM layout that is automatically updated.

→ All updates are first entered in NSM then eventually copied into DSM mirror.




A CASE FOR FRACTURED MIRRORS  
VLDB 2002



# FRACTURED MIRRORS

Store a second copy of the database in a DSM layout that is automatically updated.

→ All updates are first entered in NSM then eventually copied into DSM mirror.




A CASE FOR FRACTURED MIRRORS  
VLDB 2002



# FRACTURED MIRRORS

Store a second copy of the database in a DSM layout that is automatically updated.

→ All updates are first entered in NSM then eventually copied into DSM mirror.




A CASE FOR FRACTURED MIRRORS  
VLDB 2002



# FRACTURED MIRRORS

Store a second copy of the database in a DSM layout that is automatically updated.

→ All updates are first entered in NSM then eventually copied into DSM mirror.



A CASE FOR FRACTURED MIRRORS  
VLDB 2002

# DELTA STORE

---


Stage updates to the database in an NSM table.  
A background thread migrates updates from delta store and applies them to DSM data.



# DELTA STORE


---

Stage updates to the database in an NSM table.  
A background thread migrates updates from delta store and applies them to DSM data.



# DELTA STORE

Stage updates to the database in an NSM table.  
A background thread migrates updates from delta store and applies them to DSM data.



# SINGLE, FLEXIBLE ARCHITECTURE

---

Use a single execution engine architecture that is able to operate on both NSM and DSM data.

- Don't need to store two copies of the database.
- Don't need to sync multiple database segments.

Note that a DBMS can use the delta-store for NSM data with a single architecture.

# H<sub>2</sub>O ADAPTIVE STORAGE

---

Examine the access patterns of queries and then dynamically reconfigure the database to optimize decomposition and layout.

Copies columns into a new layout that is optimized for each query.

- Think of it like a mini fractured mirror.
- Use query compilation to speed up operations.



H2O: A HANDS-FREE ADAPTIVE STORE  
*SIGMOD 2014*



# H<sub>2</sub>O ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
 WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
 WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

A	B	C	D

# H<sub>2</sub>O ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
 WHERE C = "xxx"
```

```
SELECT AVG(B)  
FROM JoyStillSux  
WHERE C = "yyy"
```

```
SELECT SUM(A)  
FROM JoyStillSux
```

*Original Data*

A	B	C	D

# H<sub>2</sub>O ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
 WHERE C = "xxx"
```

```
SELECT AVG(B)  
FROM JoyStillSux  
WHERE C = "yyy"
```

```
SELECT SUM(A)  
FROM JoyStillSux
```

*Original Data*

A	B	C	D

# H<sub>2</sub>O ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
 WHERE C = "xxx"
```

```
SELECT AVG B  
FROM JoyStillSux  
WHERE C = "yyy"
```

```
SELECT SUM A  
FROM JoyStillSux
```

*Original Data*

A	B	C	D



*Adapted Data*

A	B	C	D

## H<sub>2</sub>O ADAPTIVE STORAGE

---

This approach is unable to handle updates to the database.

It also unable to store tuples in the same table in a different layout.

This is because they are missing the ability to categorize whether data is hot or cold...

# PELTON ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
 WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
 WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

A	B	C	D

# PELTON ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
  WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
  WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

A	B	C	D

# PELTON ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
  WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
  WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

	A	B	C	D
<i>Hot</i>				
<i>Cold</i>				



# PELTON ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
  WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
  WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

*Hot*

A	B	C	D

*Cold*


*Adapted Data*

A	B	C	D

A	B	C	D

A	B	C	D

# PELTON ADAPTIVE STORAGE

```
UPDATE JoyStillSux  
  SET B = 1234  
  WHERE C = "xxx"
```

```
SELECT AVG(B)  
  FROM JoyStillSux  
  WHERE C = "yyy"
```

```
SELECT SUM(A)  
  FROM JoyStillSux
```

*Original Data*

A	B	C	D



*Adapted Data*

A	B	C	D

A	B	C	D

# CATEGORIZING DATA

---

## **Choice #1: Manual Approach**

→ DBA specifies what tables should be stored as DSM.

## **Choice #2: Off-line Approach**

→ DBMS monitors access logs offline and then makes decision about what data to move to DSM.

## **Choice #3: On-line Approach**

→ DBMS tracks access patterns at runtime and then makes decision about what data to move to DSM.

## PARTING THOUGHTS

---

A flexible architecture that supports a hybrid storage model is the next major trend in DBMSs

This will enable relational DBMSs to support all known database workloads except for matrices in machine learning.



*JOY's DANK*  
**TIPS FOR  
PROFILING**

# MOTIVATION

---

Consider a hot program **Z** with two functions **foo** and **bar**.

How can we speed up **Z** with only a debugger ?

- Randomly pause it during execution
- Collect the function call stack

# RANDOM PAUSE METHOD

---

Consider this scenario

- Collected 10 call stack samples
- Say 6 out of the 10 samples were in **foo**

What percentage of time was spent in **foo**?

- Roughly 60% of the time was spent in **foo**
- Accuracy increases with # of samples

# AMDAHL'S LAW

---

Say we optimized **foo** to run 2 times faster  
What's the expected overall speedup ?

- **p** = percentage of time spent in optimized task
- **s** = speed up for the optimized task
- Overall speedup =  $\frac{1}{\frac{p}{s} + (1-p)}$  = 1.4 times faster



# AMDAHL'S LAW

---

Say we optimized **foo** to run 2 times faster

What's the expected overall speedup ?

→ 60% of time spent in **foo** drops in half

→ 40% of time spent in **bar** unaffected

→ **p** = percentage of time spent in optimized task

→ **s** = speed up for the optimized task

→ Overall speedup =  $\frac{1}{1 - p + \frac{p}{s}}$  = 1.4 times faster

# AMDAHL'S LAW

---

1 0.6 2 +0.4 1 1 0.6 2 +0.4 0.6 2 0.6 0.6 2 2 0.6  
 2 +0.4 1 0.6 2 +0.4 = 1.4 times faster

1 0.6 2 +0.4 1 1 0.6 2 +0.4 0.6 2 0.6 0.6 2 2 0.6  
 2 +0.4 1 0.6 2 +0.4 = 1.4 times faster

$$\frac{1}{2 + (1 - 0.6)} + (1 - 0.6) = \frac{1}{1.4} + 0.4 = 0.714 + 0.4 = 1.114$$

Say we optimized **foo** to run 2 times faster

What's the expected overall speedup ?

- 60% of time spent in **foo** drops in half
- 40% of time spent in **bar** unaffected

# PROFILING TOOLS FOR REAL

---

## **Choice #1: Valgrind**

- Heavyweight instrumentation framework with a lot of tools
- Sophisticated visualization tools

## **Choice #2: Perf**

- Lightweight tool that can record different kinds of events
- Console-oriented visualization tools

# CHOICE #1: VALGRIND

---

Instrumentation framework for building dynamic analysis tools

- **memcheck**: a memory error detector
- **callgrind**: a call-graph generating profiler

## CHOICE #1: VALGRIND

---

Instrumentation framework for building dynamic analysis tools

→ **memcheck**: a memory error detector

→ **callgrind**: a call-graph generating profiler

Using **callgrind** to profile the index test and Peloton in general:

```
$ valgrind --tool=callgrind --trace-children=yes  
./tests/index_test
```

```
$ valgrind --tool=callgrind --trace-children=yes  
./build/src/peloton -D data &> /dev/null&
```

# KCACHEGRIND

---

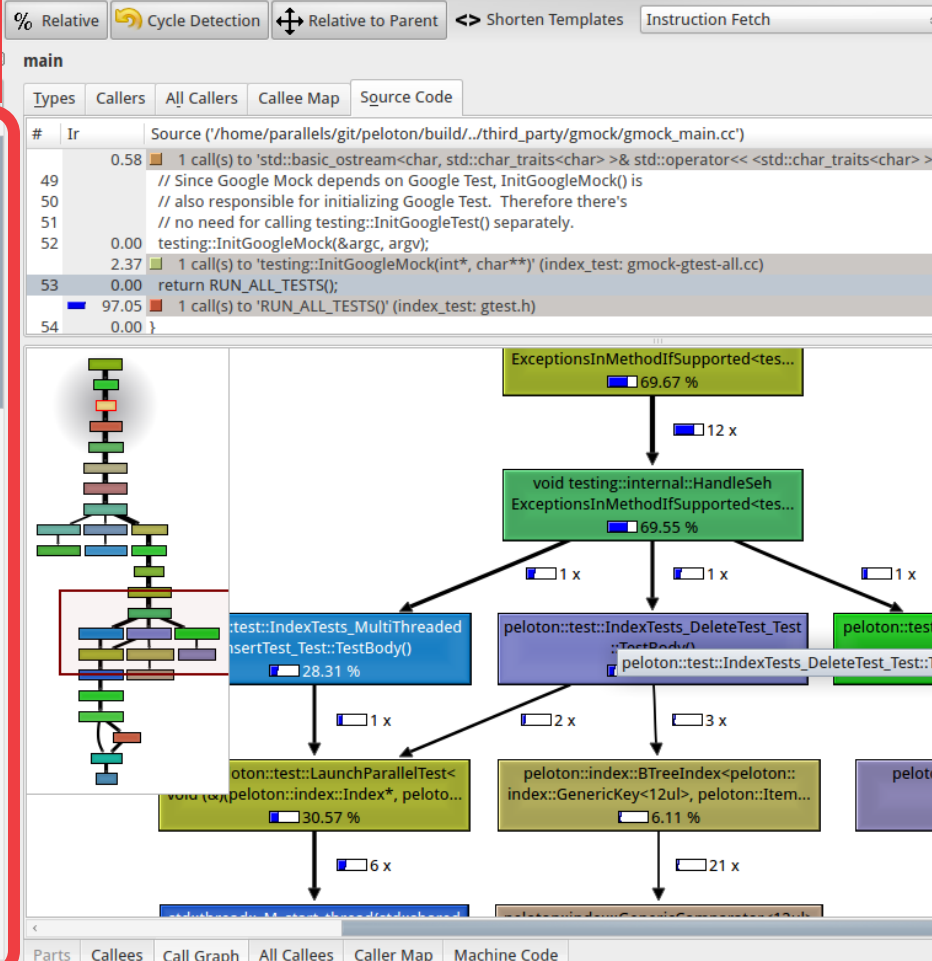
Profile data visualization tool

```
$ kcachegrind callgrind.out.12345
```



# Cumulative Time Distribution

Incl.	Self	Called	Function
92.84	0.00	(0)	0x00000000000012d0
78.31	0.01	1	_dl_start
78.30	0.01	1	_dl_sysdep_start
78.29	0.02	1	dl_main
76.43	11.20	13	_dl_relocate_object
68.77	38.98	7 312	_dl_lookup_symbol_x
29.79	23.17	7 312	do_lookup_x
11.15	0.00	1	0x00000000000406f7e
11.14	0.00	1	(below main)
7.16	0.00	4	start_thread
7.04	0.00	6	0x0000000000008d370
7.03	0.00	6	std::thread::Impl<std::Bin...
6.95	0.52	319	peloton::index::GenericCom...
6.58	0.00	1	main
6.39	0.00	1	RUN_ALL_TESTS()
6.39	0.00	1	testing::UnitTests::Run()
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	testing::internal::UnitTestIm...
5.73	3.57	7 149	check_match.9458
5.49	0.04	5	peloton::test::InsertTest(pel...
5.34	0.00	1	testing::TestCase::Run()
5.07	0.00	3	testing::TestInfo::Run()
4.93	0.04	46	peloton::index::BTreeIndex...
4.82	0.07	46	std::btree<peloton::index::G...
4.60	0.00	12	void testing::internal::Handl...
4.60	0.00	3	testing::Test::Run()
4.60	0.00	12	void testing::internal::Handl...
3.91	0.07	301	_dl_runtime_resolve
3.85	0.00	1	_libc_csu_init
3.85	0.27	301	_dl_fixup
3.69	1.06	2 187	malloc
3.59	0.53	2 024	peloton::Value::Value(pelot...
3.38	0.00	1	_dl_init
3.38	0.02	13	call_init.part.0

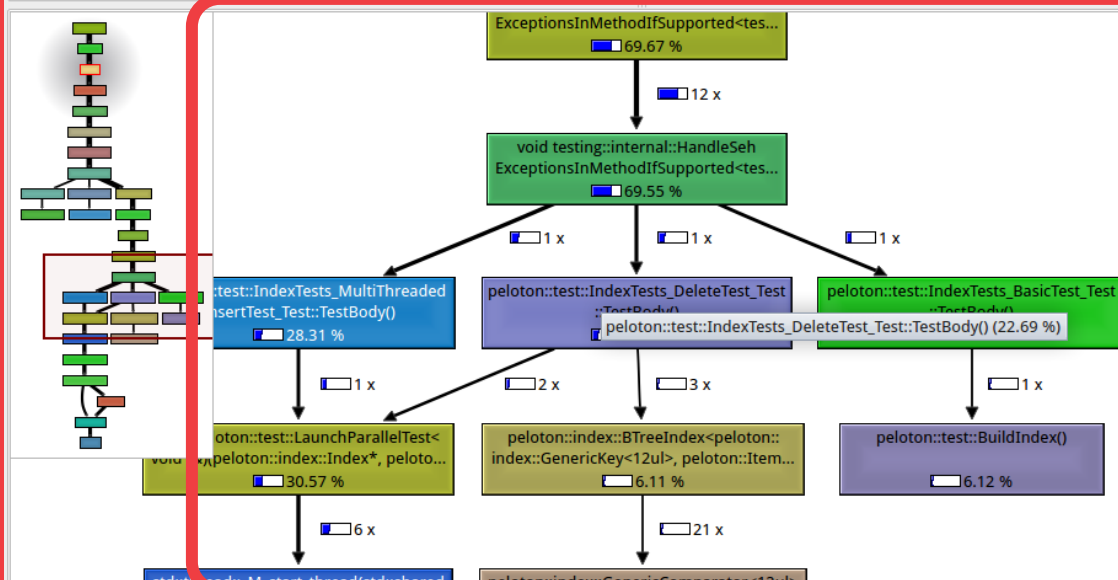




# Cumulative Time Distribution

Incl.	Self	Called	Function
92.84	0.00	(0)	0x00000000000012d0
78.31	0.01	1	_dl_start
78.30	0.01	1	_dl_sysdep_start
78.29	0.02	1	dl_main
76.43	11.20	13	_dl_relocate_object
68.77	38.98	7 312	_dl_lookup_symbol_x
29.79	23.17	7 312	do_lookup_x
11.15	0.00	1	0x00000000000406f7e
11.14	0.00	1	(below main)
7.16	0.00	4	start_thread
7.04	0.00	6	0x0000000000008d370
7.03	0.00	6	std::thread::Impl<std::Bin...
6.95	0.52	319	peloton::index::GenericCom...
6.58	0.00	1	main
6.39	0.00	1	RUN_ALL_TESTS()
6.39	0.00	1	testing::UnitTests::Run()
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	testing::internal::UnitTestIm...
5.73	3.57	7 149	check_match.9458
5.49	0.04	5	peloton::test::InsertTest(pel...
5.34	0.00	1	testing::TestCase::Run()
5.07	0.00	3	testing::TestInfo::Run()
4.93	0.04	46	peloton::index::BTreeIndex...
4.82	0.07	46	std::btree<peloton::index::G...
4.60	0.00	12	void testing::internal::Handl...
4.60	0.00	3	testing::Test::Run()
4.60	0.00	12	void testing::internal::Handl...
3.91	0.07	301	_dl_runtime_resolve
3.85	0.00	1	_libc_csu_init
3.85	0.27	301	_dl_fixup
3.69	1.06	2 187	malloc
3.59	0.53	2 024	peloton::Value::Value(pelot...
3.38	0.00	1	_dl_init
3.38	0.02	13	call_init.part.0

## Callgraph View



## CHOICE #2: PERF

---

Tool for using the performance counters subsystem in Linux.

- **-e** = sample the event cycles at the user level only
- **-c** = collect a sample every 2000 occurrences of event

```
$ perf record -e cycles:u -c 2000  
./tests/index_test
```

Uses counters for tracking events

- On counter overflow, the kernel records a sample
- Sample contains info about program execution

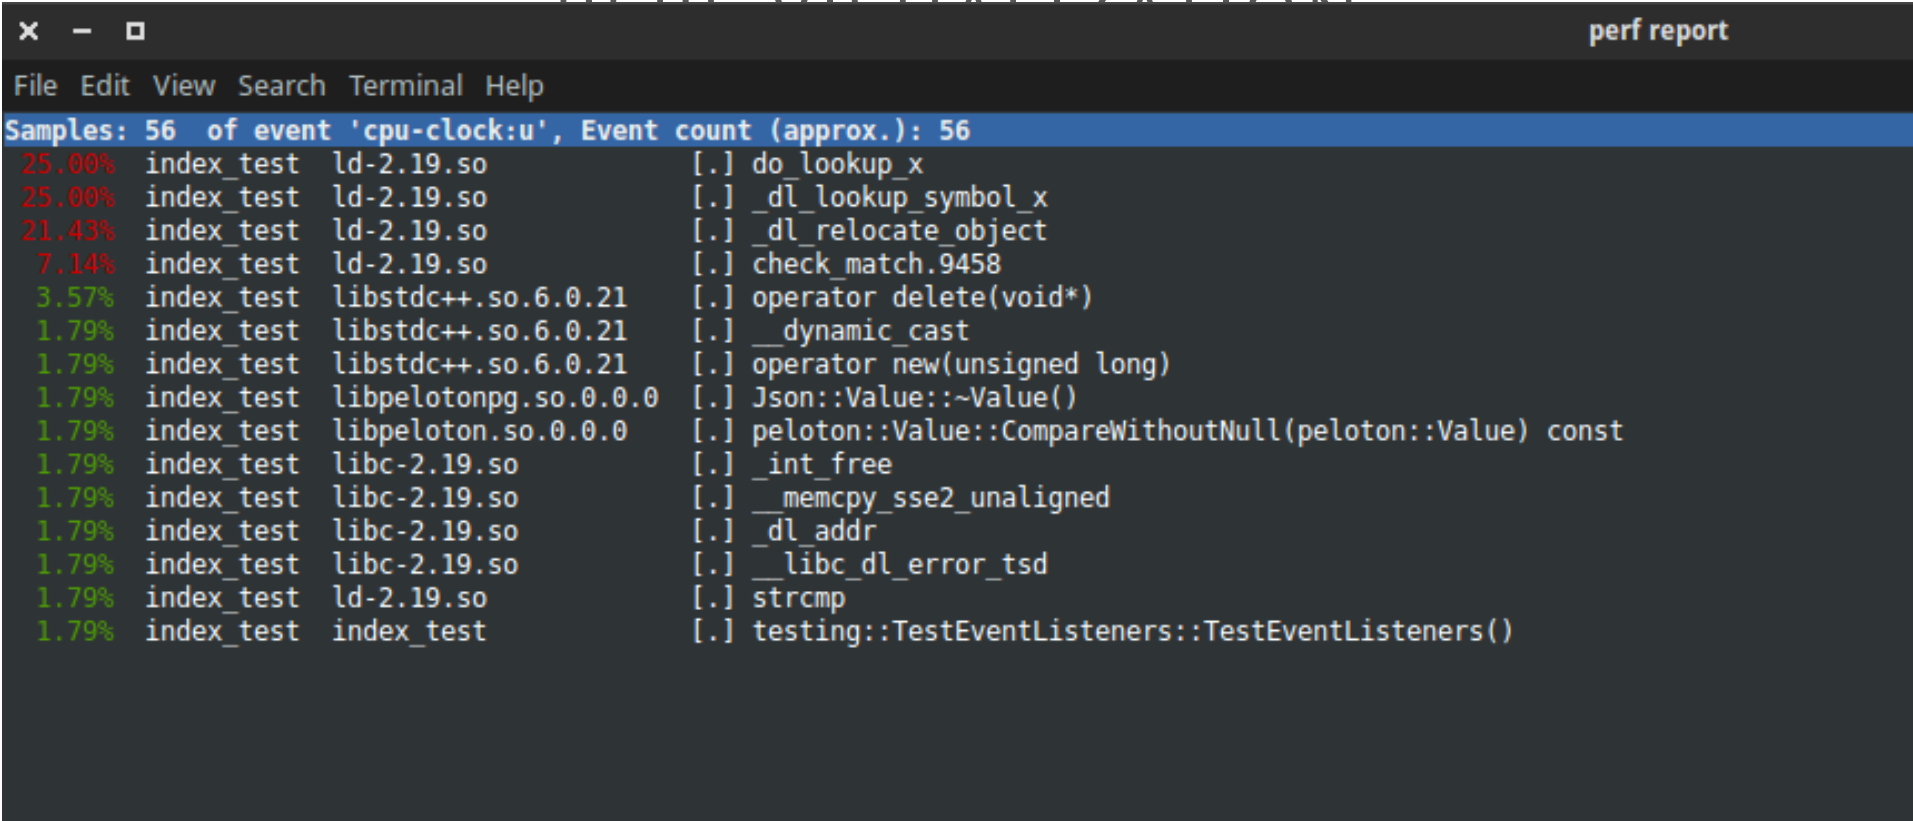
# PERF VISUALIZATION

---

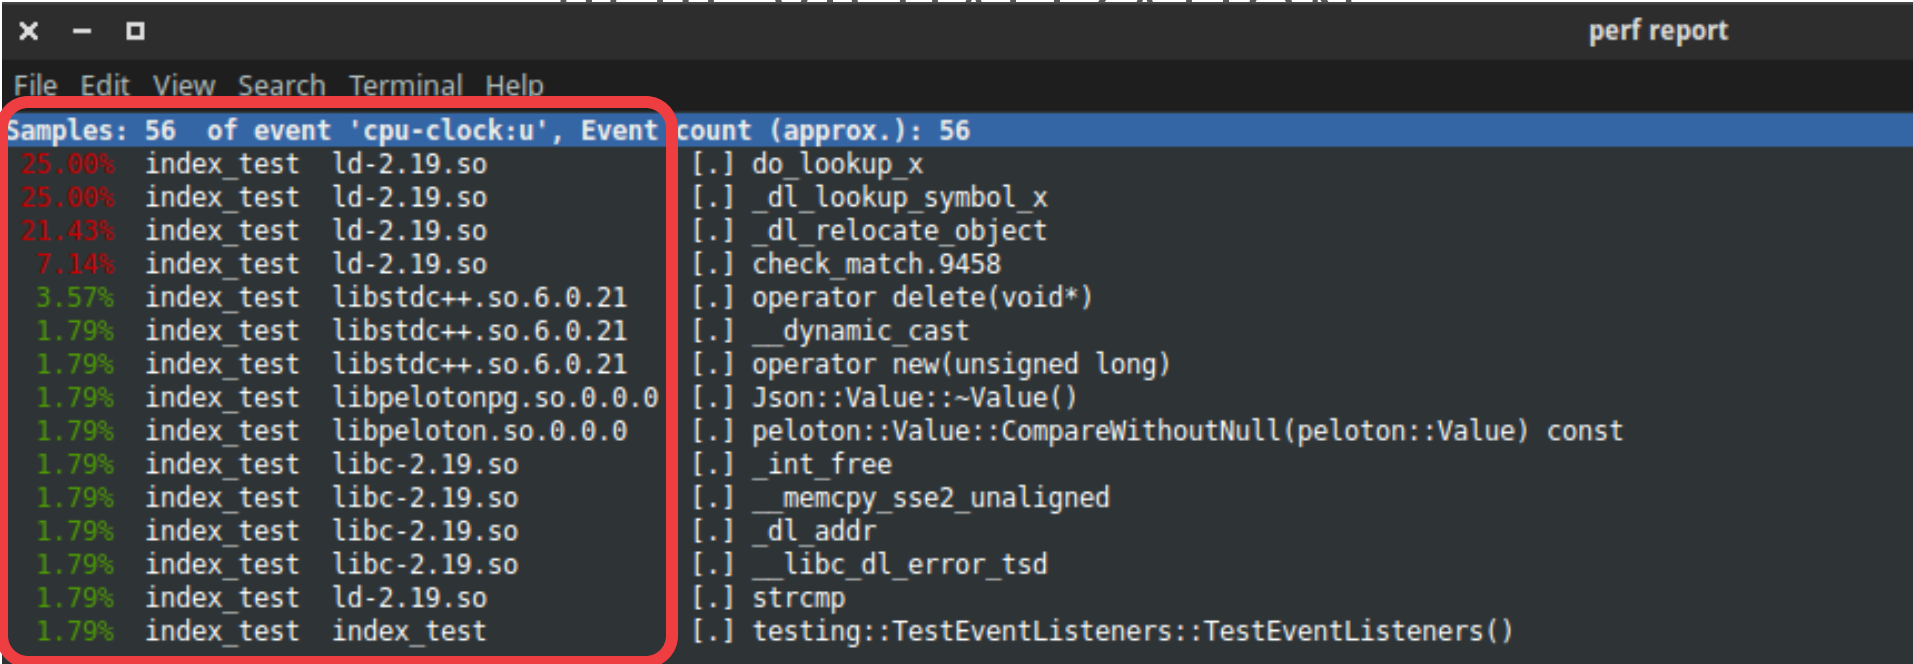
We can also use **perf** to visualize the generated profile for our application.

```
$ perf report
```

## DEDE VISUALIZATION



## DEDE VISUALIZATION



Cumulative Time  
Distribution

# PERF EVENTS

---

Supports several other events like:

- L1-dcache-load-misses
- branch-misses

To see a list of events:

```
$ perf list
```

Another usage example:

```
$ perf record -e cycles,LLC-load-misses -c 2000  
./tests/index_test
```

# REFERENCES

---

## Valgrind

- [The Valgrind Quick Start Guide](#)
- [Callgrind](#)
- [Kcachegrind](#)
- [Tips for the Profiling/Optimization process](#)

## Perf

- [Perf Tutorial](#)
- [Perf Examples](#)
- [Perf Analysis Tools](#)